

```
!pip install catboost

Collecting catboost
  Downloading catboost-1.2.8-cp312-cp312-
manylinux2014_x86_64.whl.metadata (1.2 kB)
Requirement already satisfied: graphviz in
/usr/local/lib/python3.12/dist-packages (from catboost) (0.21)
Requirement already satisfied: matplotlib in
/usr/local/lib/python3.12/dist-packages (from catboost) (3.10.0)
Requirement already satisfied: numpy<3.0,>=1.16.0 in
/usr/local/lib/python3.12/dist-packages (from catboost) (2.0.2)
Requirement already satisfied: pandas>=0.24 in
/usr/local/lib/python3.12/dist-packages (from catboost) (2.2.2)
Requirement already satisfied: scipy in
/usr/local/lib/python3.12/dist-packages (from catboost) (1.16.2)
Requirement already satisfied: plotly in
/usr/local/lib/python3.12/dist-packages (from catboost) (5.24.1)
Requirement already satisfied: six in /usr/local/lib/python3.12/dist-
packages (from catboost) (1.17.0)
Requirement already satisfied: python-dateutil>=2.8.2 in
/usr/local/lib/python3.12/dist-packages (from pandas>=0.24->catboost)
(2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in
/usr/local/lib/python3.12/dist-packages (from pandas>=0.24->catboost)
(2025.2)
Requirement already satisfied: tzdata>=2022.7 in
/usr/local/lib/python3.12/dist-packages (from pandas>=0.24->catboost)
(2025.2)
Requirement already satisfied: contourpy>=1.0.1 in
/usr/local/lib/python3.12/dist-packages (from matplotlib->catboost)
(1.3.3)
Requirement already satisfied: cyeler>=0.10 in
/usr/local/lib/python3.12/dist-packages (from matplotlib->catboost)
(0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in
/usr/local/lib/python3.12/dist-packages (from matplotlib->catboost)
(4.60.1)
Requirement already satisfied: kiwisolver>=1.3.1 in
/usr/local/lib/python3.12/dist-packages (from matplotlib->catboost)
(1.4.9)
Requirement already satisfied: packaging>=20.0 in
/usr/local/lib/python3.12/dist-packages (from matplotlib->catboost)
(25.0)
Requirement already satisfied: pillow>=8 in
/usr/local/lib/python3.12/dist-packages (from matplotlib->catboost)
(11.3.0)
Requirement already satisfied: pyparsing>=2.3.1 in
/usr/local/lib/python3.12/dist-packages (from matplotlib->catboost)
(3.2.5)
Requirement already satisfied: tenacity>=6.2.0 in
```

```
/usr/local/lib/python3.12/dist-packages (from plotly->catboost)
(8.5.0)
Downloading catboost-1.2.8-cp312-cp312-manylinux2014_x86_64.whl (99.2
MB)
----- 99.2/99.2 MB 8.2 MB/s eta
0:00:00

import pandas as pd
import numpy as np
import os
import kagglehub
from sklearn.model_selection import train_test_split
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer
from xgboost import XGBClassifier
from lightgbm import LGBMClassifier
from catboost import CatBoostClassifier
from imblearn.over_sampling import SMOTE
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score,
    roc_auc_score, classification_report, confusion_matrix,
    roc_curve, precision_recall_curve
)
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')

# --- Step 1: Load Dataset ---
path = kagglehub.dataset_download("uciml/default-of-credit-card-
clients-dataset")
print("Dataset path:", path)

Using Colab cache for faster access to the 'default-of-credit-card-
clients-dataset' dataset.
Dataset path: /kaggle/input/default-of-credit-card-clients-dataset

# Find the file (xls or csv)
file = [f for f in os.listdir(path) if f.endswith('.xls') or
f.endswith('.xlsx') or f.endswith('.csv')][0]

if file.endswith('.xls', '.xlsx'):
    df = pd.read_excel(os.path.join(path, file), header=1)
else:
    df = pd.read_csv(os.path.join(path, file))

print(df)
```

PAY_3	\	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2
0		1	200000.0	2	2	1	24	2	2
-1		2	120000.0	2	2	2	26	-1	2
0		3	90000.0	2	2	2	34	0	0
0		4	50000.0	2	2	1	37	0	0
0		5	50000.0	1	2	1	57	-1	0
-1		...	...	...	...	...	...	...	...
...		...	...	...	...	...	...	...	...
29995	29996	220000.0	1	3	1	39	0	0	0
0		29996	29997	150000.0	1	3	2	43	-1
-1		29997	29998	30000.0	1	2	2	37	4
2		29998	29999	80000.0	1	3	1	41	1
0		29999	30000	50000.0	1	2	1	46	0
0									
PAY_4	\	...	BILL_AMT4	BILL_AMT5	BILL_AMT6	PAY_AMT1	PAY_AMT2		
0		-1	...	0.0	0.0	0.0	0.0	689.0	
1		0	...	3272.0	3455.0	3261.0	0.0	1000.0	
2		0	...	14331.0	14948.0	15549.0	1518.0	1500.0	
3		0	...	28314.0	28959.0	29547.0	2000.0	2019.0	
4		0	...	20940.0	19146.0	19131.0	2000.0	36681.0	
...		...	...	...	...	...	...	...	...
29995	0	...	88004.0	31237.0	15980.0	8500.0	20000.0		
29996	-1	...	8979.0	5190.0	0.0	1837.0	3526.0		
29997	-1	...	20878.0	20582.0	19357.0	0.0	0.0		
29998	0	...	52774.0	11855.0	48944.0	85900.0	3409.0		
29999	0	...	36535.0	32428.0	15313.0	2078.0	1800.0		
PAY_AMT3	PAY_AMT4	PAY_AMT5	PAY_AMT6						

```

default.payment.next.month
0          0.0      0.0      0.0      0.0
1
1      1000.0    1000.0      0.0    2000.0
1
2      1000.0    1000.0    1000.0    5000.0
0
3      1200.0    1100.0    1069.0    1000.0
0
4     10000.0    9000.0     689.0     679.0
0
...
...
29995   5003.0    3047.0    5000.0    1000.0
0
29996   8998.0     129.0      0.0      0.0
0
29997  22000.0    4200.0    2000.0    3100.0
1
29998   1178.0    1926.0    52964.0    1804.0
1
29999   1430.0     1000.0    1000.0    1000.0
1

[30000 rows x 25 columns]

# Clean column names
df.columns = df.columns.str.strip().str.lower().str.replace(" ", "_").str.replace(".", "_")

# Drop ID column if exists
if 'id' in df.columns:
    df.drop(columns=['id'], inplace=True)

print("Columns:", df.columns.tolist())

Columns: ['limit_bal', 'sex', 'education', 'marriage', 'age', 'pay_0',
'pay_2', 'pay_3', 'pay_4', 'pay_5', 'pay_6', 'bill_amt1', 'bill_amt2',
'bill_amt3', 'bill_amt4', 'bill_amt5', 'bill_amt6', 'pay_amt1',
'pay_amt2', 'pay_amt3', 'pay_amt4', 'pay_amt5', 'pay_amt6',
'default_payment_next_month']

# --- Step 2: Split features and target ---
target_col = 'default_payment_next_month'
X = df.drop(columns=[target_col])
y = df[target_col]

# --- Step 3: Preprocessing pipeline ---
num_cols = X.select_dtypes(include=[np.number]).columns.tolist()

```

```

num_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())
])

preprocessor = ColumnTransformer([
    ('num', num_pipeline, num_cols)
])

X_processed = preprocessor.fit_transform(X)

print(X)
print(y)
print(X_processed)

```

	limit_bal	sex	education	marriage	age	pay_0	pay_2	pay_3
pay_4	20000.0	2	2	1	24	2	2	-1
0	120000.0	2	2	2	26	-1	2	0
-1	90000.0	2	2	2	34	0	0	0
0	50000.0	2	2	1	37	0	0	0
0	50000.0	1	2	1	57	-1	0	-1
0	50000.0	1	2	1	57	-1	0	-1
...	...	...	...	...	...	...	...	...
...	...	...	...	...	...	...	...	...
29995	220000.0	1	3	1	39	0	0	0
0	150000.0	1	3	2	43	-1	-1	-1
29996	150000.0	1	3	2	43	-1	-1	-1
-1	30000.0	1	2	2	37	4	3	2
29997	30000.0	1	2	2	37	4	3	2
-1	80000.0	1	3	1	41	1	-1	0
29998	80000.0	1	3	1	41	1	-1	0
0	50000.0	1	2	1	46	0	0	0
29999	50000.0	1	2	1	46	0	0	0
0	50000.0	1	2	1	46	0	0	0
pay_amt1	pay_5	...	bill_amt3	bill_amt4	bill_amt5	bill_amt6		
0	-2	...	689.0	0.0	0.0	0.0		
0.0	0	...	2682.0	3272.0	3455.0	3261.0		
1	0	...	13559.0	14331.0	14948.0	15549.0		
0.0	0	...	49291.0	28314.0	28959.0	29547.0		
1518.0	0	...	49291.0	28314.0	28959.0	29547.0		
2000.0	0	...	49291.0	28314.0	28959.0	29547.0		

```

4          0 ...    35835.0    20940.0    19146.0    19131.0
2000.0
...
29995      0 ...    208365.0   88004.0    31237.0    15980.0
8500.0
29996      0 ...     3502.0     8979.0     5190.0     0.0
1837.0
29997      0 ...    2758.0    20878.0    20582.0    19357.0
0.0
29998      0 ...    76304.0   52774.0    11855.0    48944.0
85900.0
29999      0 ...    49764.0    36535.0    32428.0    15313.0
2078.0

          pay_amt2  pay_amt3  pay_amt4  pay_amt5  pay_amt6
0        689.0      0.0      0.0      0.0      0.0
1       1000.0    1000.0    1000.0      0.0    2000.0
2       1500.0    1000.0    1000.0   1000.0    5000.0
3       2019.0    1200.0    1100.0   1069.0    1000.0
4      36681.0   10000.0    9000.0     689.0     679.0
...
29995    20000.0    5003.0    3047.0    5000.0    1000.0
29996    3526.0    8998.0     129.0      0.0      0.0
29997      0.0   22000.0    4200.0    2000.0    3100.0
29998    3409.0    1178.0    1926.0   52964.0    1804.0
29999    1800.0    1430.0    1000.0    1000.0    1000.0

[30000 rows x 23 columns]
0      1
1      1
2      0
3      0
4      0
...
29995    0
29996    0
29997    1
29998    1
29999    1
Name: default_payment_next_month, Length: 30000, dtype: int64
[[-1.13672015  0.81016074  0.18582826 ... -0.30806256 -0.31413612
 -0.29338206]
 [-0.3659805   0.81016074  0.18582826 ... -0.24422965 -0.31413612
 -0.18087821]
 [-0.59720239  0.81016074  0.18582826 ... -0.24422965 -0.24868274
 -0.01212243]
...
[-1.05964618 -1.23432296  0.18582826 ... -0.03996431 -0.18322937]
```

```

-0.11900109]
[-0.67427636 -1.23432296  1.45111372 ... -0.18512036  3.15253642
-0.19190359]
[-0.90549825 -1.23432296  0.18582826 ... -0.24422965 -0.24868274
-0.23713013]]
```

# --- Step 4: Train-test split ---

```
X_train, X_test, y_train, y_test = train_test_split(
    X_processed, y, test_size=0.2, random_state=42, stratify=y
)
```

# --- Step 5: Apply SMOTE to training data ---

```
smote = SMOTE(random_state=42)
X_train_resampled, y_train_resampled = smote.fit_resample(X_train,
y_train)
```

```
print(f"Before SMOTE: {np.bincount(y_train)}")
print(f"After SMOTE: {np.bincount(y_train_resampled)}")
```

```
Before SMOTE: [18691  5309]
After SMOTE: [18691 18691]
```

# --- Step 6: Define evaluation function ---

```
def evaluate_model(y_test, y_pred, y_proba, model_name):
    acc = accuracy_score(y_test, y_pred)
    prec = precision_score(y_test, y_pred)
    rec = recall_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)
    auc = roc_auc_score(y_test, y_proba)

    print(f"\n{model_name} Performance:")
    print(f"Accuracy: {acc:.4f}")
    print(f"Precision: {prec:.4f}")
    print(f"Recall: {rec:.4f}")
    print(f"F1-score: {f1:.4f}")
    print(f"AUC: {auc:.4f}")
    print(classification_report(y_test, y_pred))
```

# Confusion Matrix

```
plt.figure(figsize=(5, 4))
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, fmt='d',
cmap='Blues')
plt.title(f'{model_name} - Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.tight_layout()
plt.show()
```

# ROC Curve

```
fpr, tpr, _ = roc_curve(y_test, y_proba)
```

```

plt.figure(figsize=(5, 4))
plt.plot(fpr, tpr, label=f'AUC = {auc:.3f}')
plt.plot([0, 1], [0, 1], '--', color='gray')
plt.title(f'{model_name} - ROC Curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.tight_layout()
plt.show()

# Precision-Recall Curve
prec_curve, rec_curve, _ = precision_recall_curve(y_test, y_proba)
plt.figure(figsize=(5, 4))
plt.plot(rec_curve, prec_curve)
plt.title(f'{model_name} - Precision-Recall Curve')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.tight_layout()
plt.show()

# --- Step 7: Train and evaluate XGBoost ---
xgb_model = XGBClassifier(use_label_encoder=False,
                           eval_metric='logloss', random_state=42)
xgb_model.fit(X_train_resampled, y_train_resampled)
xgb_pred = xgb_model.predict(X_test)
xgb_proba = xgb_model.predict_proba(X_test)[:, 1]
evaluate_model(y_test, xgb_pred, xgb_proba, "XGBoost")

```

XGBoost Performance:

Accuracy: 0.8002

Precision: 0.5653

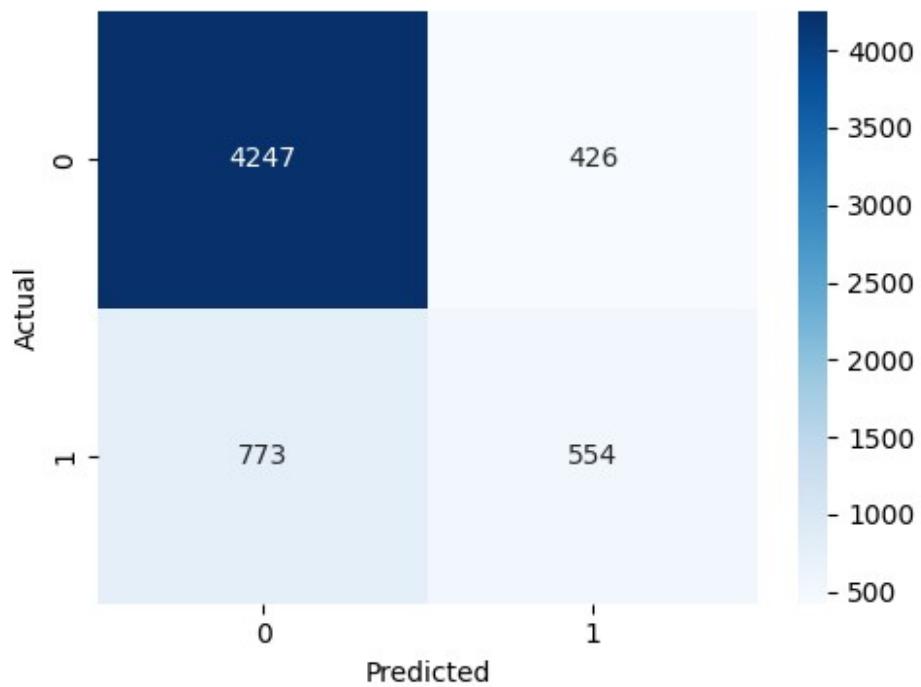
Recall: 0.4175

F1-score: 0.4803

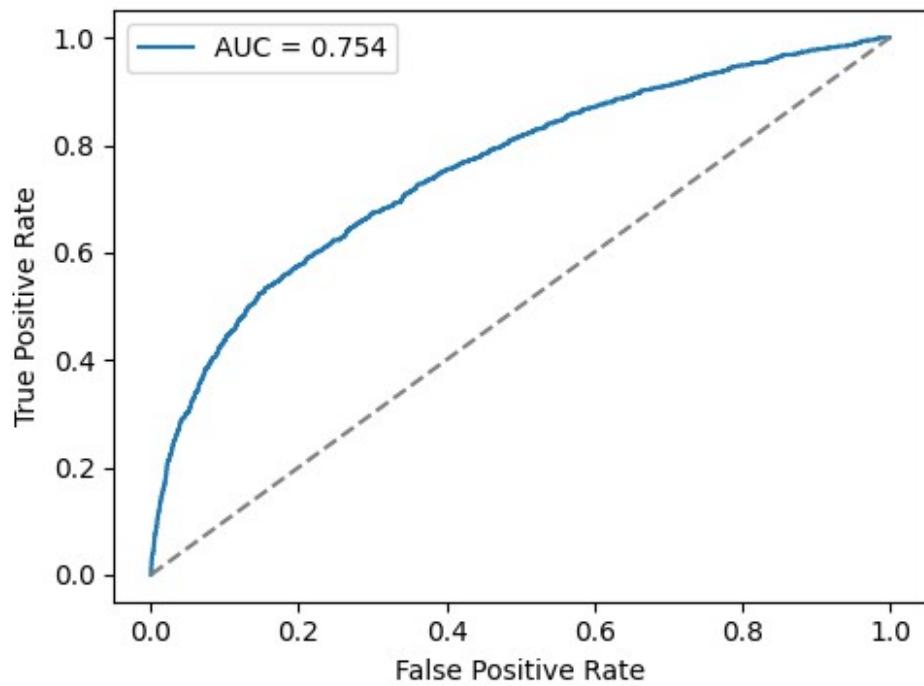
AUC: 0.7539

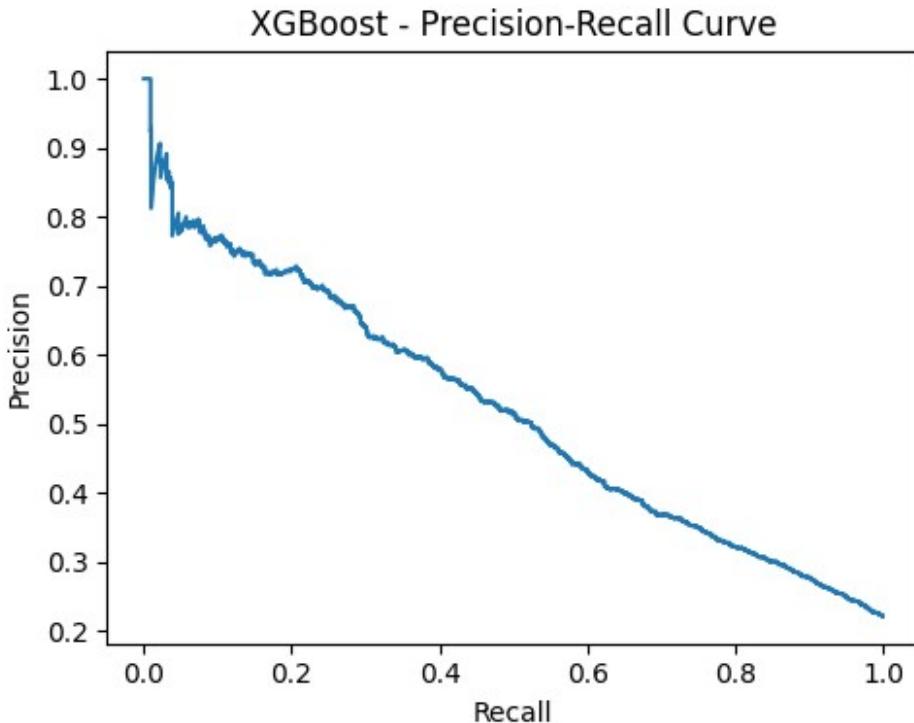
	precision	recall	f1-score	support
0	0.85	0.91	0.88	4673
1	0.57	0.42	0.48	1327
accuracy			0.80	6000
macro avg	0.71	0.66	0.68	6000
weighted avg	0.78	0.80	0.79	6000

XGBoost - Confusion Matrix



XGBoost - ROC Curve





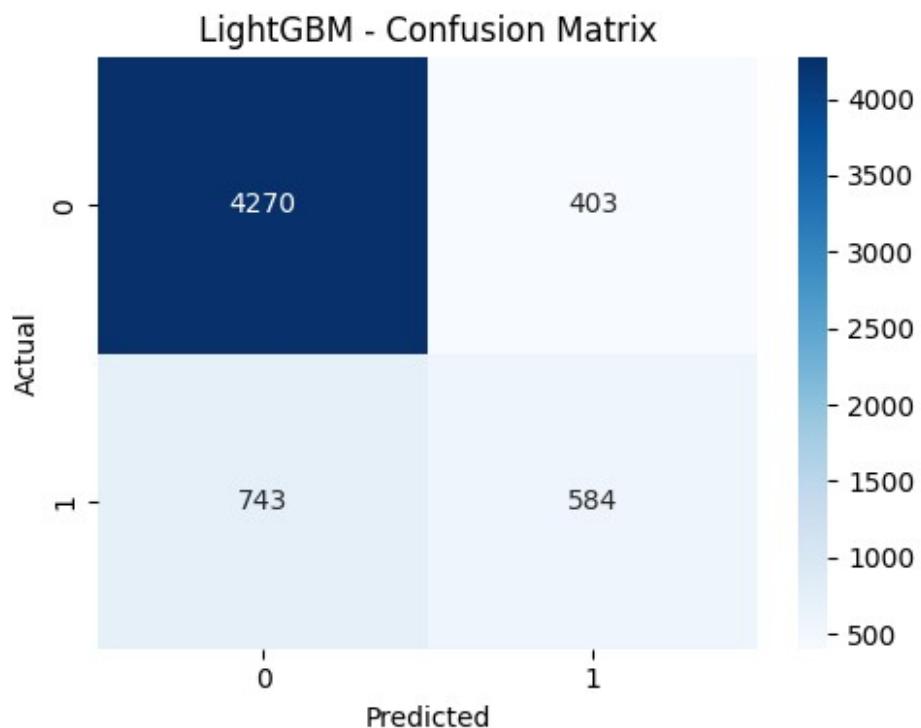
```
# --- Step 8: Train and evaluate LightGBM ---
from lightgbm import LGBMClassifier

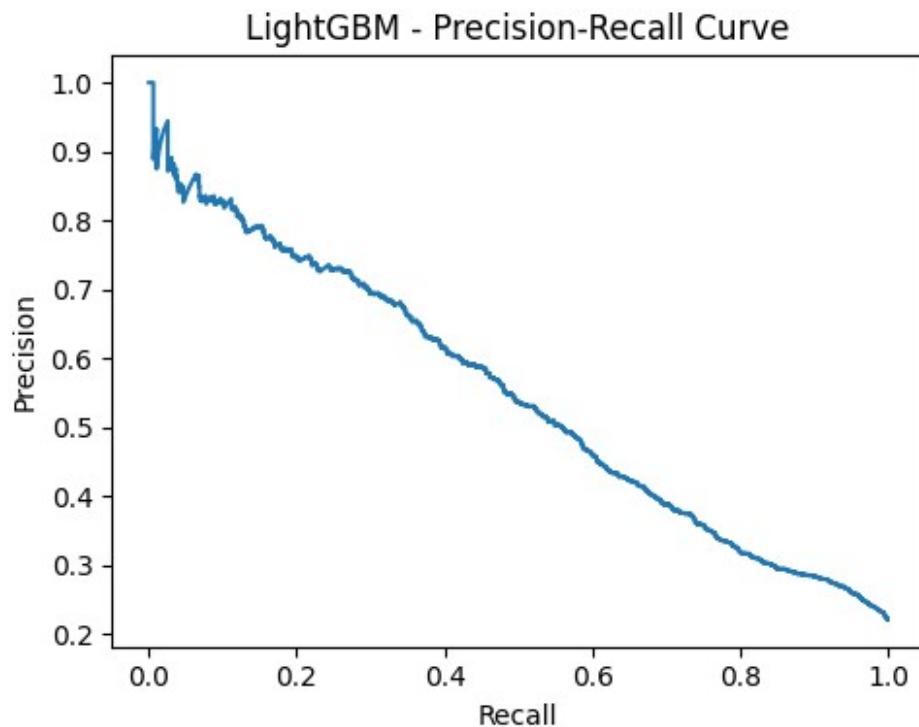
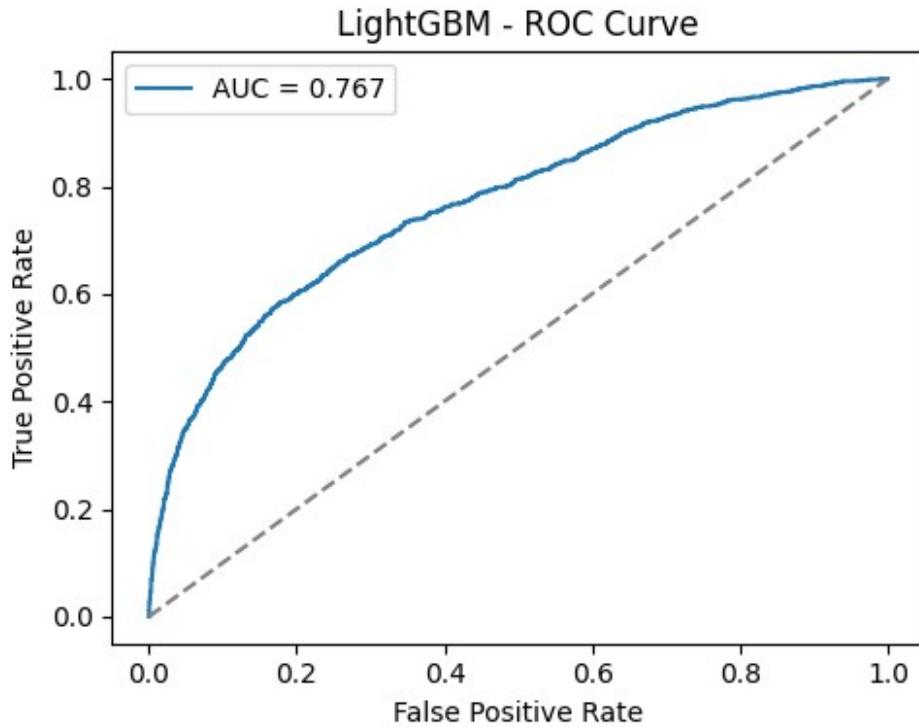
lgbm_model = LGBMClassifier(random_state=42)
lgbm_model.fit(X_train_resampled, y_train_resampled)
lgbm_pred = lgbm_model.predict(X_test)
lgbm_proba = lgbm_model.predict_proba(X_test)[:, 1]
evaluate_model(y_test, lgbm_pred, lgbm_proba, "LightGBM")

[LightGBM] [Info] Number of positive: 18691, number of negative: 18691
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead
of testing was 0.003322 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true` .
[LightGBM] [Info] Total Bins 5599
[LightGBM] [Info] Number of data points in the train set: 37382,
number of used features: 23
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 ->
initscore=0.000000

LightGBM Performance:
Accuracy: 0.8090
Precision: 0.5917
Recall: 0.4401
F1-score: 0.5048
AUC: 0.7674
      precision    recall   f1-score   support
```

	0	0.85	0.91	0.88	4673
	1	0.59	0.44	0.50	1327
accuracy				0.81	6000
macro avg		0.72	0.68	0.69	6000
weighted avg		0.79	0.81	0.80	6000





```
# --- Step 9: Train and evaluate CatBoost ---  
from catboost import CatBoostClassifier  
  
cat_model = CatBoostClassifier(verbose=0, random_state=42)
```

```

cat_model.fit(X_train_resampled, y_train_resampled)
cat_pred = cat_model.predict(X_test)
cat_proba = cat_model.predict_proba(X_test)[:, 1]
evaluate_model(y_test, cat_pred, cat_proba, "CatBoost")

```

CatBoost Performance:

Accuracy: 0.8125

Precision: 0.6194

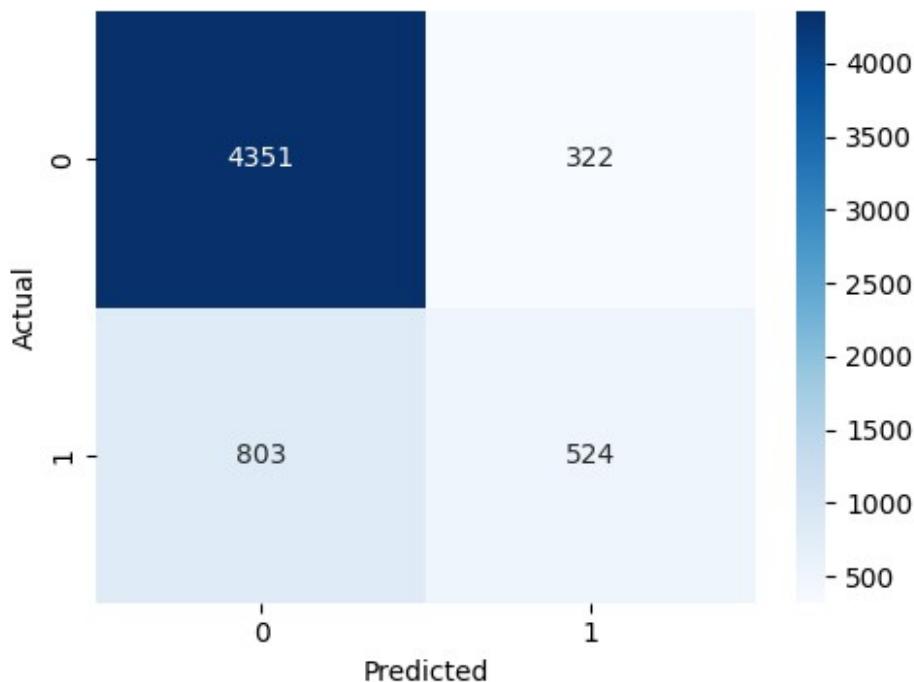
Recall: 0.3949

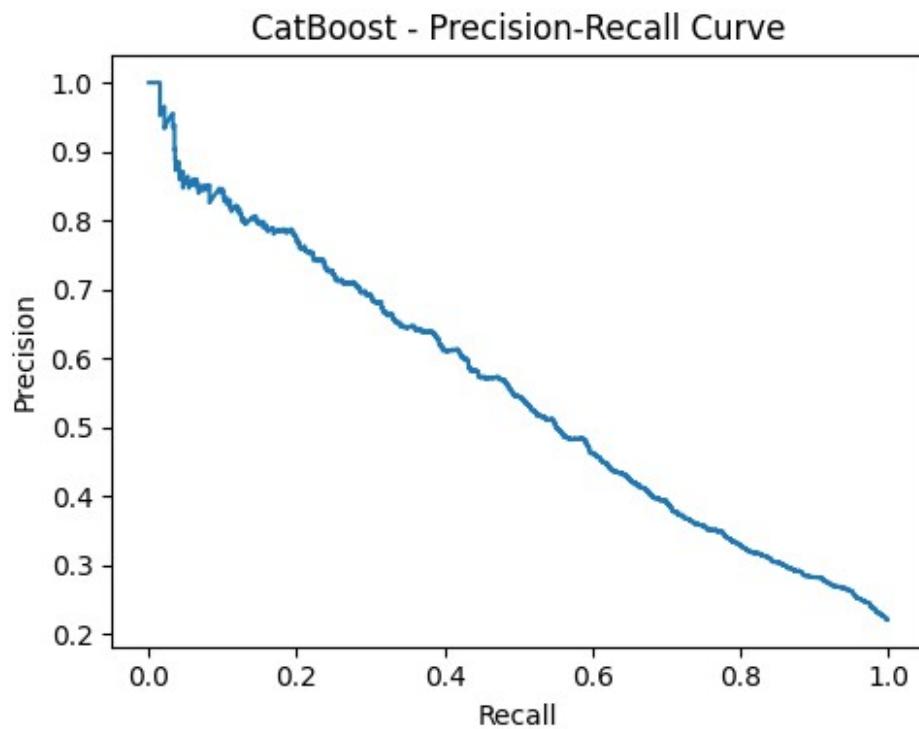
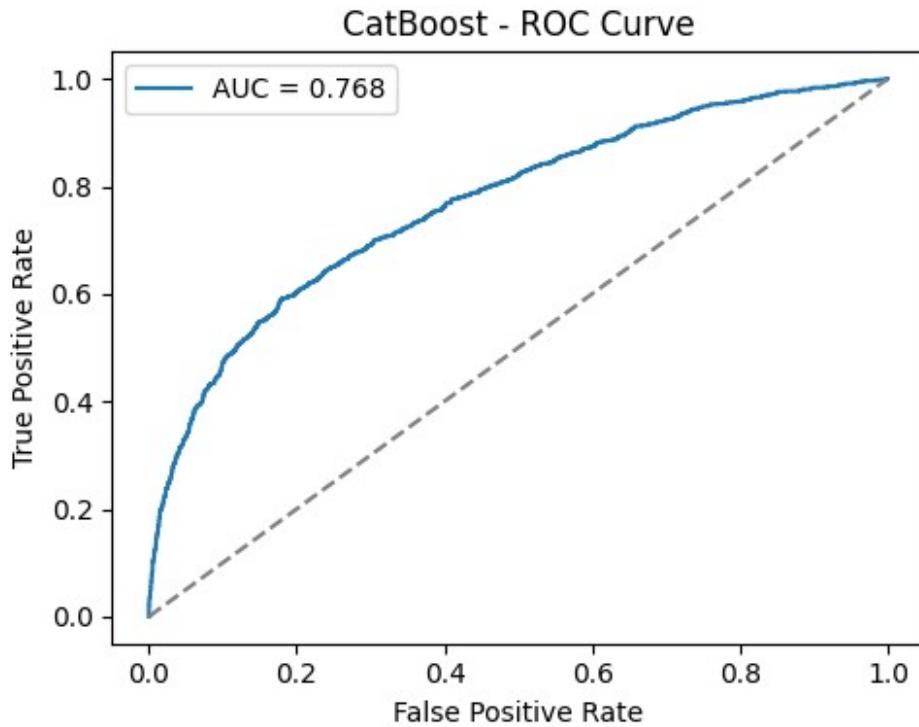
F1-score: 0.4823

AUC: 0.7682

	precision	recall	f1-score	support
0	0.84	0.93	0.89	4673
1	0.62	0.39	0.48	1327
accuracy			0.81	6000
macro avg	0.73	0.66	0.68	6000
weighted avg	0.79	0.81	0.80	6000

CatBoost - Confusion Matrix





```
# --- Step 10: Feature importance plots ---
def plot_feature_importance(model, feature_names, model_name):
    importances = model.feature_importances_
    imp_df = pd.DataFrame({
```

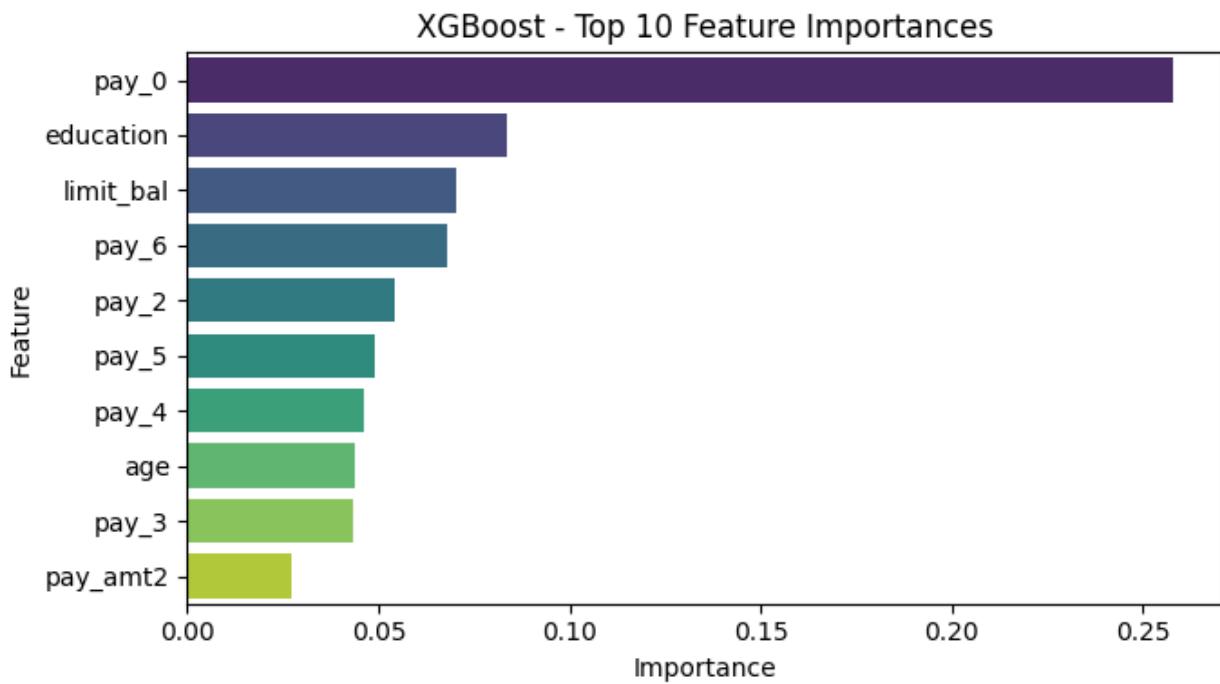
```

        'Feature': feature_names,
        'Importance': importances
    }).sort_values('Importance', ascending=False).head(10)

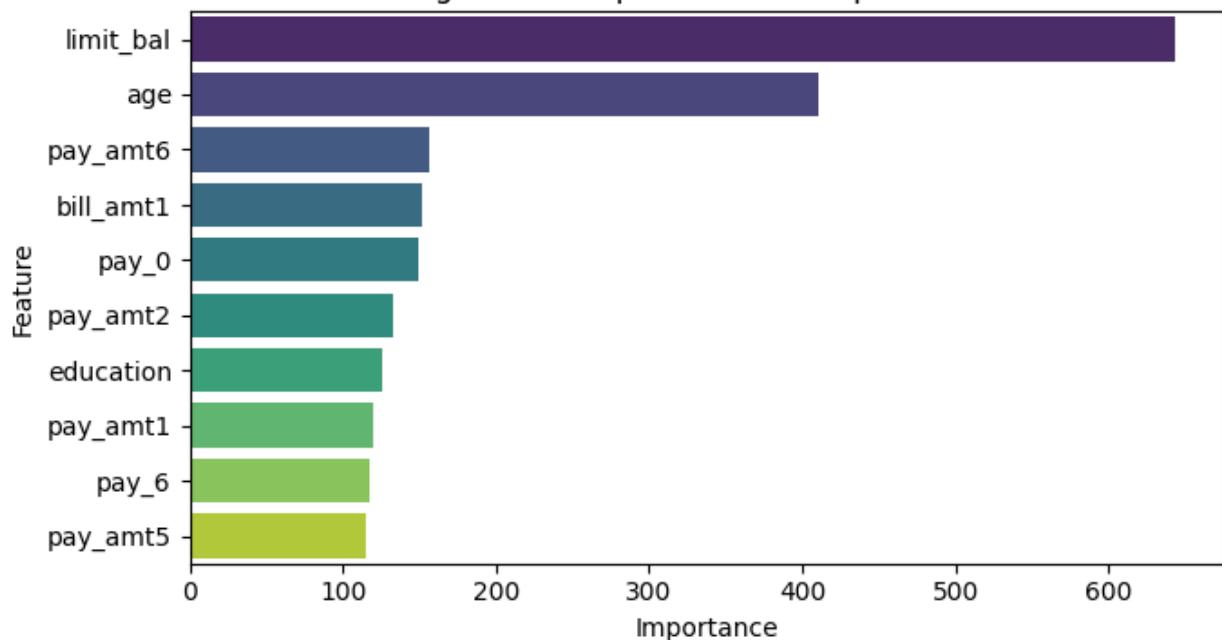
    plt.figure(figsize=(7, 4))
    sns.barplot(x='Importance', y='Feature', data=imp_df,
    palette='viridis')
    plt.title(f'{model_name} - Top 10 Feature Importances')
    plt.tight_layout()
    plt.show()

# Original feature names after preprocessing (num_cols)
plot_feature_importance(xgb_model, num_cols, "XGBoost")
plot_feature_importance(lgbm_model, num_cols, "LightGBM")
plot_feature_importance(cat_model, num_cols, "CatBoost")

```



LightGBM - Top 10 Feature Importances



CatBoost - Top 10 Feature Importances

